

DD2424: Assignment 2

Sam Shahriari

April 18, 2024

1 Gradient

To check that my analytical calculation of the gradient was correct, I compared it to a numerical estimation for the same points and features. The numerical estimation was calculated by the given function `ComputeGradsNumSlow()` which uses centered difference formula. The batches tested were all possible combinations of `featureSize = 20`, `batchSize $\in \{1, 10, 100\}$` and `$\lambda \in \{0, .1, 1\}$` . None of the absolute or relative errors were above 10^{-6} and therefore I draw the conclusion that my implementation is correct. More detailed results of the testing can be found in [Appendix A](#).

Without any regularization, it is very easy to overfit the model. This can be seen in [Figure 1](#).




Figure 1: Overfitting the model

2 Cyclical learning

A cyclical learning scheme means that we continuously change one of the hyperparameters from a min value to a max value and then back to a min value. Here we used the cyclical approach on the learning rate which changed between $\eta_{min} = 10^{-5}$ and $\eta_{max} = 10^{-1}$. The effect of varying the learning rate can be seen in [Figure 2](#) and even more clearly in [Figure 3](#).

When examining the graph it can be seen that the y-values mostly goes in the right direction but in an area close the highest learning rate they get worse. This is probably because when using a high lambda, the step might be too big so the minima is jumped over. This could be a good attribute if the loss function contains saddle points or local minimas as they then hopefully will be skipped.

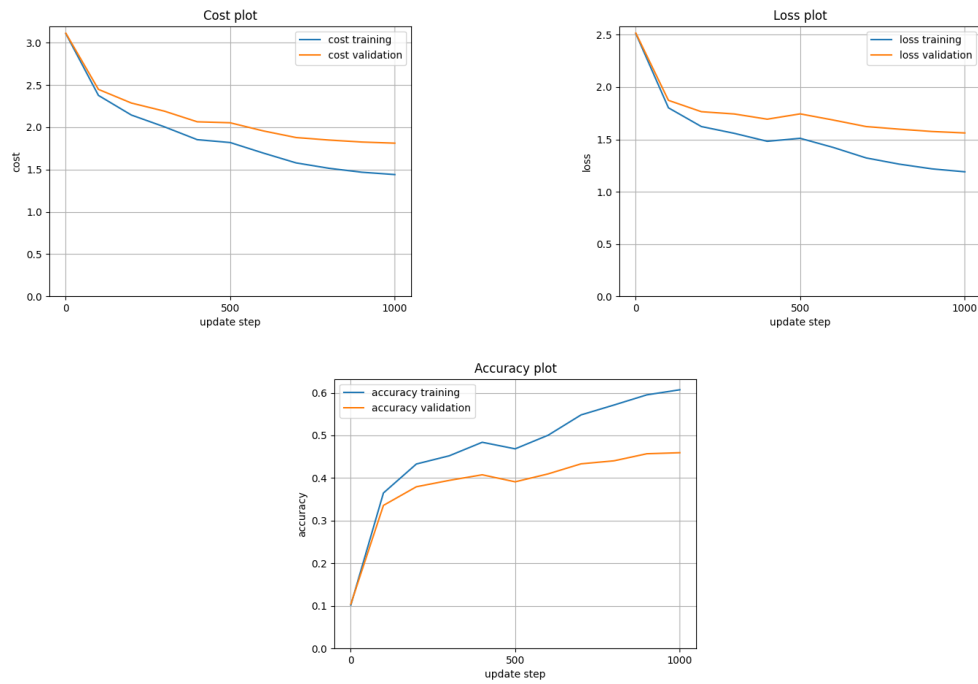


Figure 2: Plots for one cycle with $n_s = 500$

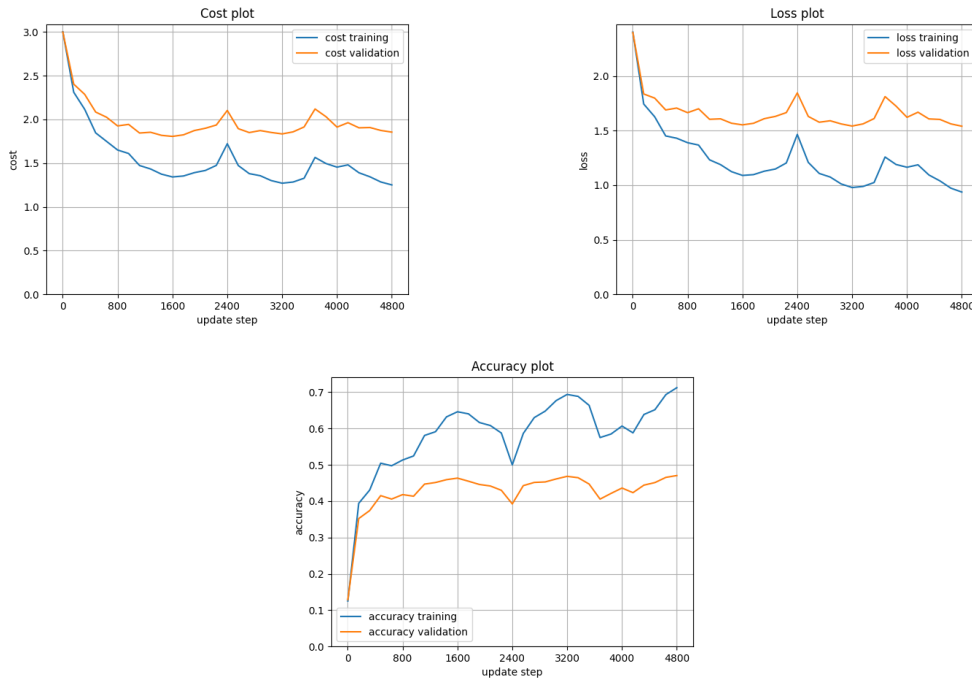


Figure 3: Plots for three cycles with $n_s = 800$

3 Search for lambda

To search for a good lambda value, the training set was extended. It now contained all training batches except for 5000 images that were used as the validation set.

The search was a coarse search with a wide possible range of lambda. 8 random lambda values were chosen in the \log_{10} range from -1 to -5 . Other parameter settings were cycles = 2, batch size = 100, $n_s = 900$, $\eta_{min} = 10^{-5}$, $\eta_{max} = 10^{-1}$. The three best performing lambdas were:

- Accuracy 0.532 lambda 7.807727344534329e-05 log lambda -4.107475361175347
- Accuracy 0.5296 lambda 5.066346576621312e-05 log lambda -4.295305104485099
- Accuracy 0.5294 lambda 7.193692167309888e-05 log lambda -4.143048150462331

The best lambda was then used for a finer search. Now 20 values were generated in the log range from -4.107 ± 1 . The other hyperparameters were the same as in the coarse search. The three best performing lambdas were:

- Accuracy 0.531 lambda $9.867193573189926e-06$ log lambda -5.005806351787182
- Accuracy 0.529 lambda $1.6403528938636582e-05$ log lambda -4.785062710872495
- Accuracy 0.5272 lambda $2.3557420875260744e-05$ log lambda -4.627872258917016

Lastly, the best lambda = $9.867e-06$ was used to train the model. When testing the model on the test data an accuracy of 51.37% was achieved. A plot of the loss function can be seen in [Figure 4](#) and all the lambdas and their respective accuracy can be seen in [Appendix B](#).

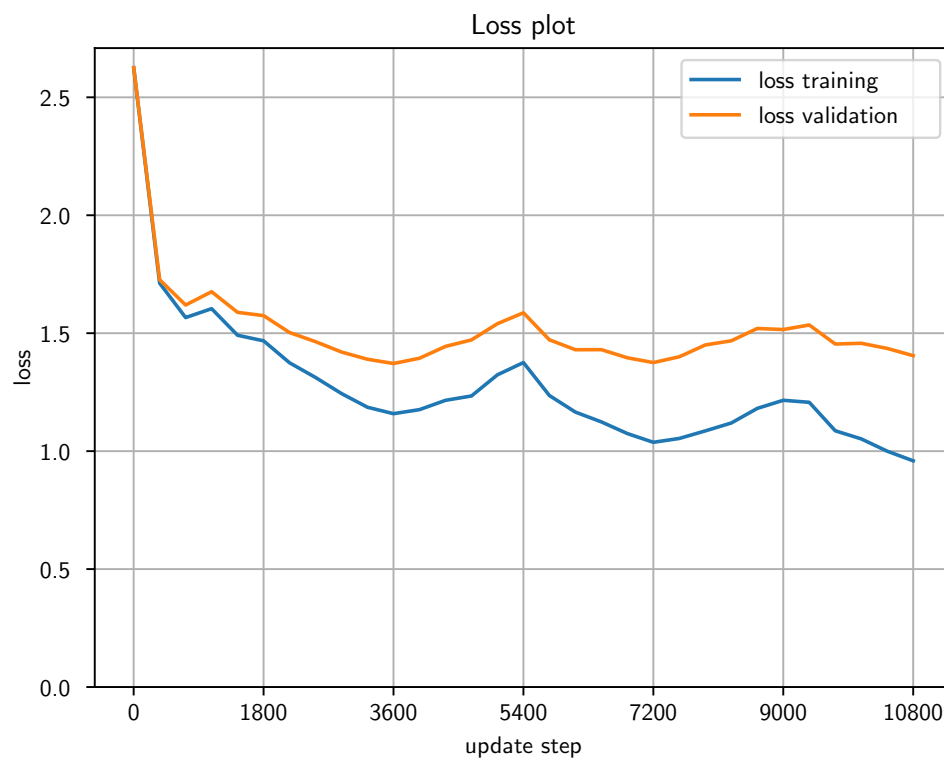


Figure 4: Plot of the loss function for the training and validation set with the best lambda.

A Gradients

```
num features: 20 , batch size: 1 , lambda: 0
SLOW W abs wrong 0 biggest error 2.4892828250078214e-10
SLOW W rel wrong 0 biggest error 1.4533802739282618e-08
SLOW b abs wrong 0 biggest error 6.958188192296433e-11
SLOW b rel wrong 0 biggest error 3.4129318763622067e-10
-----
num features: 20 , batch size: 1 , lambda: 0.1
SLOW W abs wrong 0 biggest error 4.316231330681042e-10
SLOW W rel wrong 0 biggest error 2.7179142571849796e-09
SLOW b abs wrong 0 biggest error 6.958188192296433e-11
SLOW b rel wrong 0 biggest error 3.4129318763622067e-10
-----
num features: 20 , batch size: 1 , lambda: 1
SLOW W abs wrong 0 biggest error 1.4254555591453055e-09
SLOW W rel wrong 0 biggest error 9.170683667909907e-10
SLOW b abs wrong 0 biggest error 6.415323966502129e-10
SLOW b rel wrong 0 biggest error 3.1466616103752504e-09
-----
num features: 20 , batch size: 10 , lambda: 0
SLOW W abs wrong 0 biggest error 4.117283293816887e-10
SLOW W rel wrong 0 biggest error 2.4653246494206414e-08
SLOW b abs wrong 0 biggest error 3.1131446032173216e-10
SLOW b rel wrong 0 biggest error 1.4792496788051695e-09
-----
num features: 20 , batch size: 10 , lambda: 0.1
SLOW W abs wrong 0 biggest error 5.233947858451771e-10
SLOW W rel wrong 0 biggest error 3.4908850003620592e-09
SLOW b abs wrong 0 biggest error 3.1131446032173216e-10
SLOW b rel wrong 0 biggest error 1.4792496788051695e-09
-----
num features: 20 , batch size: 10 , lambda: 1
SLOW W abs wrong 0 biggest error 1.805463070714275e-09
SLOW W rel wrong 0 biggest error 1.1683430450338027e-09
SLOW b abs wrong 0 biggest error 8.080397739806955e-10
SLOW b rel wrong 0 biggest error 3.83950226689584e-09
-----
num features: 20 , batch size: 100 , lambda: 0
SLOW W abs wrong 0 biggest error 5.002005569725715e-10
SLOW W rel wrong 0 biggest error 3.516490864551905e-08
```

```

SLOW b abs wrong 0 biggest error 2.42636927910711e-10
SLOW b rel wrong 0 biggest error 1.9957499952210056e-09
-----
num features: 20 , batch size: 100 , lambda: 0.1
SLOW W abs wrong 0 biggest error 5.088013725618179e-10
SLOW W rel wrong 0 biggest error 3.3926608475673607e-09
SLOW b abs wrong 0 biggest error 2.42636927910711e-10
SLOW b rel wrong 0 biggest error 1.9957499952210056e-09
-----
num features: 20 , batch size: 100 , lambda: 1
SLOW W abs wrong 0 biggest error 1.6478069464476164e-09
SLOW W rel wrong 0 biggest error 1.0661789264572265e-09
SLOW b abs wrong 0 biggest error 7.109335689592378e-10
SLOW b rel wrong 0 biggest error 5.8476081075959645e-09

```

B Lambda Random Search

```

Accuracy 0.532 lambda 7.807727344534329e-05 log lambda -4.107475361175347
Accuracy 0.5296 lambda 5.066346576621312e-05 log lambda -4.295305104485099
Accuracy 0.5294 lambda 7.193692167309888e-05 log lambda -4.143048150462331
Accuracy 0.5236 lambda 0.00014498005602996882 log lambda -3.8386917367433813
Accuracy 0.5224 lambda 0.00010803538625357892 log lambda -3.966433971017319
Accuracy 0.5224 lambda 3.197602887656392e-05 log lambda -4.495175472549843
Accuracy 0.522 lambda 0.0016228962075828747 log lambda -2.7897092546133395
Accuracy 0.516 lambda 2.931955461229069e-05 log lambda -4.532842631264662
-----
Accuracy 0.531 lambda 9.867193573189926e-06 log lambda -5.005806351787182
Accuracy 0.529 lambda 1.6403528938636582e-05 log lambda -4.785062710872495
Accuracy 0.5272 lambda 2.3557420875260744e-05 log lambda -4.627872258917016
Accuracy 0.5262 lambda 0.00010832213245853509 log lambda -3.965282798937192
Accuracy 0.5256 lambda 9.069945553883424e-06 log lambda -5.042395319965149
Accuracy 0.5244 lambda 7.7774484176332e-05 log lambda -4.109162860580757
Accuracy 0.5234 lambda 1.089283904004913e-05 log lambda -4.962858913748786
Accuracy 0.5218 lambda 0.0004735627803922964 log lambda -3.324622438264533
Accuracy 0.5216 lambda 0.000607347932639988 log lambda -3.2165624424743715
Accuracy 0.52 lambda 1.828514484964466e-05 log lambda -4.737901594956661
Accuracy 0.5194 lambda 3.353723742505652e-05 log lambda -4.474472714552444
Accuracy 0.5194 lambda 2.917295194356437e-05 log lambda -4.535019623400148
Accuracy 0.5186 lambda 0.0004622448632289073 log lambda -3.335127906309916
Accuracy 0.5174 lambda 7.063692063734561e-05 log lambda -4.150968241754363

```

```
Accuracy 0.5168 lambda 0.0005695846990260709 log lambda -3.244441685835034
Accuracy 0.5156 lambda 4.6161234167673446e-05 log lambda -4.335722588391209
Accuracy 0.5146 lambda 0.00010063522872220608 log lambda -3.997249962009483
Accuracy 0.5136 lambda 1.2155852830676228e-05 log lambda -4.915214566506879
Accuracy 0.5114 lambda 5.137162709657223e-05 log lambda -4.289276678623103
Accuracy 0.5072 lambda 0.0007257027866074966 log lambda -3.1392412092790263
-----
```

```
Accuracy on test data: 0.5137
```